


Jacob Chalk



Research Focus

My research focuses on Multimodal Learning, aiming to develop methods that integrate complementary sources of information to better understand dynamic real-world environments. My previous work has explored multimodal approaches to egocentric video understanding, including audio-visual learning, predicting upcoming object interactions using eye-gaze and 3D information, and long-term 3D dynamic object tracking.

Current Role

University of Bristol

Senior Research Associate
Researching Multimodal Learning.

Bristol, UK

September 2026 – Current

Previous Roles

University of Bristol

Research Associate
Conducted research on 4D Video Understanding.

Bristol, UK

January 2026 – September 2026

NAVER LABS Europe

Industrial Internship - Visual Representation Learning Team
Carried out research on long-term 3D multi-object tracking, supervised by **Dr. Diane Larlus**.

Grenoble, France

February 2025 – January 2026

Academic Qualifications

University of Bristol

PhD in Computer Vision - **Examiners: Prof. Hilde Kuehne and Prof. Andrew Owens**
Thesis Title: Leveraging Multimodal Data for Egocentric Video Understanding, supervised by **Prof. Dima Damen**.

Bristol, UK

September 2021 – September 2025

University of Bristol

MEng in Computer Science - **First Class Honours**
Dissertation Title: Video GANs for Human-Object Interactions, supervised by **Prof. Dima Damen**.

Bristol, UK

September 2017 – September 2021

Teaching

University of Bristol

Teaching Assistant
Software Product Engineering (2nd Year Module), Games Project (3rd Year Module), Computer Graphics (3rd Year Module)
Image Processing and Computer Vision (3rd Year Module), Applied Deep Learning (4th Year Module)

Bristol, UK

September 2019 – September 2024

Publications

- **J. Chalk**, S. Sinha, D. Damen, Y. Kalantidis, and D. Larlus, “Whareformer: Learning to Track What is Where in Long Egocentric Videos” in **ECCV 2026**. [\[Page\]](#) | [\[Paper\]](#)
- M. Hatano*, S. Sinha*, **J. Chalk**, W. Li, H. Saito, and D. Damen, “Prime and Reach: Synthesising Body Motion for Gaze-Primed Object Reach” in **arXiv preprint arXiv:2512.16456, 2025**. [\[Page\]](#) | [\[Paper\]](#)
- T. Perrett*, A. Darkhalil*, S. Sinha*, O. Emara*, S. Pollard*, K. Parida*, K. Liu*, P. Gatti*, S. Bansal*, K. Flanagan*, **J. Chalk***, Z. Zhu*, R. Guerrier*, F. Abdelazim*, B. Zhu, D. Moltisanti, M. Wray, H. Doughty, and D. Damen, “HD-EPIC: A Highly-Detailed Egocentric Video Dataset” in **CVPR 2025**. [\[Page\]](#) | [\[Paper\]](#)
- C. Plizzari, S. Goel, T. Perrett, **J. Chalk**, A Kanazawa, D Damen, “Spatial Cognition from Egocentric Video: Out of Sight, Not Out of Mind” in **3DV 2025**. [\[Page\]](#) | [\[Paper\]](#)
- **J. Chalk***, J. Huh*, E. Kazakos, A. Zisserman, and D. Damen, “TIM: A Time Interval Machine for Audio-Visual Action Recognition” in **CVPR 2024**. [\[Page\]](#) | [\[Paper\]](#)
- J. Huh*, **J. Chalk***, E. Kazakos, D. Damen, and A. Zisserman, “EPIC-SOUNDS: A Large-Scale Dataset of Actions that Sound” in **ICASSP 2023 & TPAMI 2025 (extended version)**. [\[Page\]](#) | [\[Paper\]](#)

*: Equal Contribution

Datasets

HD-EPIC

Highly Detailed Egocentric Video Dataset

Released February 2025

Paper Published in CVPR 2025

HD-EPIC is an extensive dataset manually annotated with highly detailed and interconnected ground-truth labels covering recipe steps, fine-grained actions, ingredients with nutritional values, moving objects, and audio annotations. All annotations are grounded in 3D through digital twinning of the scene, fixtures, object locations, and primed with gaze. Footage is collected from unscripted recordings in diverse home environments, making HD-EPIC the first dataset collected in-the-wild but with detailed annotations matching those in controlled lab environments. The dataset consists of 41 hours of video in 9 kitchens with digital twins of 413 kitchen fixtures, capturing 69 recipes, 59K fine-grained actions, 51K audio events, 20K object movements and 37K object masks lifted to 3D. On average, there are 263 annotations per minute across the unscripted videos. HD-EPIC also presents a challenging VQA benchmark of 26K questions assessing the capability to recognise recipes, ingredients, nutrition, fine-grained actions, 3D perception, object motion, and gaze direction.

Roles:

- Gathered audio annotations matching a similar process to EPIC-Sounds (see below)
- Computed and visualised eye-gaze priming statistics for all 3D objects start and end locations. This involved projecting the eye-gaze direction ray into the 3D environment and measuring when the ray intersected with the objects 3D-bounding box before it was picked-up, or the end locations bounding box when placed down.
- Created eye-gaze related VQA questions, including “*What is the camera wearer currently looking at?*”, which involved finding the intersection between the projected gaze ray and 3D environment, as well as “*What will the person interact with next?*” which involved using objects whose start locations were primed with eye-gaze before being picked up to indicate it was about to be moved by the camera wearer.
- Ran LLM inference of Llama 3.2 (Vision-Instruct-90B) on the VQA questions

EPIC-Sounds

Audio-Centric Annotated Dataset

Released January 2023

Paper Published in ICASSP 2023 & TPAMI 2025

EPIC-Sounds is an audio-only dataset, gathered across 100 hours of untrimmed audios from the videos of EPIC-KITCHENS. This dataset was motivated by the observation that audio-visual networks tend to use both temporally and semantically identical annotations for both modalities, yet audio and video can significantly differ even when describing the same action.

Roles:

- Post-processed the labelled segments from annotators, this included:
 - Manually correcting typos in the free-form text descriptions.
 - Removing erroneous/redundant annotations.
 - Categorising the free-form descriptions into the 44 classes in the dataset.
 - Decomposing challenging segments containing multiple overlapping sounds into distinct annotations.
 - Training binary classifiers on the head classes of the dataset to clean the training set.
- Trained and ran quantitative analysis on the baseline models.
- Visualisation of the class distribution for the dataset.
- Ran analyses on the gathered data, such as the duration distribution of all classes and interplay between the visual and audio labels.
- Set and currently manage two challenges on the dataset: Audio-Based Interaction Recognition and Audio-Based Interaction Detection

Honours and Awards

Outstanding Reviewer

European Conference on Computer Vision

Awarded in recognition of meaningful peer review contributions.

August 2026

ECVA Conference (ECCV)

EgoVis 2024/25 Distinguished Paper Awards

Third Joint Egocentric Vision Workshop

Awarded to our 2025 CVPR paper: “HD-EPIC: A Highly-Detailed Egocentric Video Dataset.”

June 2026

IEEE/CVF Conference (CVPR)

Outstanding Reviewer

International Conference on Computer Vision

Awarded in recognition of meaningful peer review contributions.

October 2025

IEEE/CVF Conference (ICCV)

Outstanding Reviewer

May 2025

Conference on Computer Vision and Pattern Recognition
Awarded in recognition of meaningful peer review contributions.

IEEE/CVF Conference (CVPR)

EgoVis 2022/23 Distinguished Paper Awards

June 2024

First Joint Egocentric Vision Workshop

IEEE/CVF Conference (CVPR)

Awarded to our 2023 ICASSP paper: "EPIC-Sounds: A large-scale dataset of actions that sound."

EPIC-KITCHENS Challenges Winner

June 2024

First Joint Egocentric Vision Workshop

IEEE/CVF Conference (CVPR)

Audio-Based Interaction Recognition (2nd), Audio-Based Interaction Detection (2nd), Action Detection (3rd)

Top 5 Third Year MEng Computer Science/Computer Science with Maths Student

October 2020

Netcraft

University of Bristol

Awarded to the Top-5 performing Computer Science and Computer Science with Maths students.

Research Activities

Workshops

- UniWorld Workshop organiser (ECCV 2026).
- EPIC-KITCHENS Challenges & HD-EPIC Challenge organiser for Third Joint Egocentric Vision Workshop (CVPR 2026).
- 2nd Workshop on World Models communication chair (ICLR 2026).
- EPIC-KITCHENS Challenges organiser for the First Joint Egocentric Vision Workshop (CVPR 2024).

Presentations

- Whareformer: Learning to Track What is Where in Long Egocentric Videos. *Presented as part of Professor Dima Damen's talk, "Modelling in an Ego-Sensed World" at the [Third Workshop on 3D Human Understanding \(ECCV 2026\)](#).*
- 4D Modelling using Aria Glasses: From HD-EPIC to Whareformer. *Presented at the [Fifth Hands-on Egocentric Research Tutorial \(ECCV 2026\)](#).*
- EPIC-KITCHENS Challenges & HD-EPIC Challenge. *Presented at the [Third Joint Egocentric Vision Workshop \(CVPR 2026\)](#).*
- TIM: A Time-Interval Machine for Audio-Visual Action Recognition. *Presented at the Sight & Sound Workshop (CVPR 2024) and [Twelve Labs Multimodal Weekly #56](#).*
- EPIC-KITCHENS Challenges. *Presented at the First Joint Egocentric Vision Workshop (CVPR 2024).*
- Oral Presentation for EPIC-Sounds. *Presented at ICASSP 2023.*

Reviewing

- 2026: CVPR, ECCV, ICPR, BMVC
- 2025: CVPR, NeurIPS, ICCV, IJCV, OJSP
- 2024: ECCV, TPAMI

Technical Knowledge and Programming

Programming Languages

- Python (Strongest)
- C++ (Intermediate)
- JavaScript (Intermediate)

Frameworks

- PyTorch (Strongest)
- TensorFlow (Intermediate)

Technical Fields

- Machine Learning, particularly Deep Learning
- Computer Vision
- Data Science